

# Code Optimization: Assignment N°5 Report

Duy Le - Cédric Léonard, 6/12/2021

## Introduction

The assignment is about understanding the compiler-generated assembly language

## Question N°1

a. Compile the code with optimization level -O2 and generate assembly language code for the saxpy procedure:

The assembly code is generated by the following command: `gcc -c -g -O2 saxpy.c -std=c99 objdump -d -S saxpy.o`

Here is the assembly code generated:

```
0000000000000000 <saxpy>:
/*
  Single-precision A*X plus Y, Y = alpha*X+Y
*/

void saxpy(int n, float alpha, float *X, float *Y) {
  for (int i=0; i<n; i++)
    0:  31 c0                xor    %eax,%eax
    2:  85 ff                test   %edi,%edi
    4:  7e 25                jle    2b <saxpy+0x2b>
    6:  66 2e 0f 1f 84 00 00  nopw   %cs:0x0(%rax,%rax,1)
    d:  00 00 00
      Y[i] = alpha*X[i] + Y[i];
   10:  f3 0f 10 0c 86          movss  (%rsi,%rax,4),%xmm1
   15:  f3 0f 59 c8             mulss  %xmm0,%xmm1
   19:  f3 0f 58 0c 82          addss  (%rdx,%rax,4),%xmm1
   1e:  f3 0f 11 0c 82          movss  %xmm1, (%rdx,%rax,4)
   23:  48 83 c0 01             add     $0x1,%rax
  for (int i=0; i<n; i++)
   27:  39 c7                  cmp     %eax,%edi
   29:  7f e5                  jg      10 <saxpy+0x10>
   2b:  f3 c3                  repz retq
}
```

- How are the parameters passed to the procedure The parameters are passed to the procedure as follow: n to register EDI, alpha to RDI, \*X to RSI, \*Y to RDX
- How is the loop implemented?

```

    for (int i=0; i<n; i++)
    0:   31 c0                                xor    %eax,%eax          # Set
i = 0
    2:   85 ff                                test   %edi,%edi         # if
n = 0
    4:   7e 25                                jle    2b <saxpy+0x2b>    #
then jump to 2b: terminate the program
    # Execute other code.
    23:  48 83 c0 01                         add     $0x1,%rax         #
Increase i

    for (int i=0; i<n; i++)
    27:  39 c7                                cmp     %eax,%edi        #
Compare i and n
    29:  7f e5                                jg      10 <saxpy+0x10>   #
Jump to the arithmetic operation if the the result is n>i
    2b:  f3 c3                                repz retq                #
Otherwise, return

```

- How are the arithmetic operations done?

```

Y[i] = alpha*X[i] + Y[i];
    10:  f3 0f 10 0c 86                        movss   (%rsi,%rax,4),%xmm1  `#
xmm1 = alpha
    15:  f3 0f 59 c8                        mulss   %xmm0,%xmm1         #
xmm1 = x[i] x alpha
    19:  f3 0f 58 0c 82                        addss   (%rdx,%rax,4),%xmm1  #
xmm1 += Y[i]
    1e:  f3 0f 11 0c 82                        movss   %xmm1, (%rdx,%rax,4)  #
Y[i] = xmm1

```

- Instructions that have no effect, i.e. that are inserted by the compiler to align branch targets? 6: 66 2e 0f 1f 84 00 00 nopw  
%cs:0x0(%rax,%rax,1) # Used to align

b. Compile the saxpy procedure with full optimizations (-O3) and look at the generated assembly language code.

The assembly code is generated by the following command: `gcc -c -g saxpy.c -o saxpyo3.o -O3 --std=c99 objdump -d -S saxpyo3.o`

Here is the assembly code generated:

```
saxpyo3.o:      file format elf64-x86-64

Disassembly of section .text:

0000000000000000 <saxpy>:
/*
  Single-precision A*X plus Y, Y = alpha*X+Y
*/

void saxpy(int n, float alpha, float *X, float *Y) {
  for (int i=0; i<n; i++)
    0:  85 ff                test   %edi,%edi
    2:  0f 8e fb 00 00 00    jle    103 <saxpy+0x103>
    8:  48 8d 42 10            lea    0x10(%rdx),%rax
    c:  48 39 c6              cmp    %rax,%rsi
    f:  48 8d 46 10            lea    0x10(%rsi),%rax
   13:  0f 93 c1              setae  %cl
   16:  48 39 c2              cmp    %rax,%rdx
   19:  0f 93 c0              setae  %al
   1c:  08 c1                or     %al,%cl
   1e:  0f 84 bc 00 00 00    je     e0 <saxpy+0xe0>
   24:  83 ff 05              cmp    $0x5,%edi
   27:  0f 86 b3 00 00 00    jbe    e0 <saxpy+0xe0>
   2d:  0f 28 e0              movaps %xmm0,%xmm4
   30:  41 89 f9              mov    %edi,%r9d
   33:  31 c0                xor    %eax,%eax
   35:  41 c1 e9 02            shr    $0x2,%r9d
      Y[i] = alpha*X[i] + Y[i];
   39:  0f 57 db              xorps  %xmm3,%xmm3
   3c:  0f c6 e4 00            shufps $0x0,%xmm4,%xmm4
   40:  46 8d 04 8d 00 00 00    lea    0x0(,%r9,4),%r8d
   47:  00
  for (int i=0; i<n; i++)
   48:  31 c9                xor    %ecx,%ecx
      Y[i] = alpha*X[i] + Y[i];
   4a:  0f 28 cb              movaps %xmm3,%xmm1
   4d:  0f 28 d3              movaps %xmm3,%xmm2
   50:  83 c1 01              add    $0x1,%ecx
   53:  0f 12 0c 06            movlps (%rsi,%rax,1),%xmm1
```

```

57: 0f 12 14 02      movlps (%rdx,%rax,1),%xmm2
5b: 0f 16 4c 06 08   movhps 0x8(%rsi,%rax,1),%xmm1
60: 0f 16 54 02 08   movhps 0x8(%rdx,%rax,1),%xmm2
65: 0f 59 cc         mulps  %xmm4,%xmm1
68: 0f 58 ca         addps  %xmm2,%xmm1
6b: 0f 13 0c 02      movlps %xmm1, (%rdx,%rax,1)
6f: 0f 17 4c 02 08   movhps %xmm1,0x8(%rdx,%rax,1)
74: 48 83 c0 10      add    $0x10,%rax
78: 44 39 c9         cmp    %r9d,%ecx
7b: 72 cd         jb     4a <saxpy+0x4a>
7d: 44 39 c7         cmp    %r8d,%edi
80: 0f 84 7d 00 00 00 je     103 <saxpy+0x103>
86: 49 63 c8         movslq %r8d,%rcx
89: f3 0f 10 0c 8e   movss  (%rsi,%rcx,4),%xmm1
8e: 48 8d 04 8a      lea    (%rdx,%rcx,4),%rax
for (int i=0; i<n; i++)
92: 41 8d 48 01      lea    0x1(%r8),%ecx
    Y[i] = alpha*X[i] + Y[i];
96: f3 0f 59 c8      mulss  %xmm0,%xmm1
for (int i=0; i<n; i++)
9a: 39 cf         cmp    %ecx,%edi
    Y[i] = alpha*X[i] + Y[i];
9c: f3 0f 58 08      addss  (%rax),%xmm1
a0: f3 0f 11 08      movss  %xmm1, (%rax)
for (int i=0; i<n; i++)
a4: 7e 5d         jle    103 <saxpy+0x103>
    Y[i] = alpha*X[i] + Y[i];
a6: 48 63 c9         movslq %ecx,%rcx
for (int i=0; i<n; i++)
a9: 41 83 c0 02      add    $0x2,%r8d
    Y[i] = alpha*X[i] + Y[i];
ad: f3 0f 10 0c 8e   movss  (%rsi,%rcx,4),%xmm1
b2: 48 8d 04 8a      lea    (%rdx,%rcx,4),%rax
for (int i=0; i<n; i++)
b6: 44 39 c7         cmp    %r8d,%edi
    Y[i] = alpha*X[i] + Y[i];
b9: f3 0f 59 c8      mulss  %xmm0,%xmm1
bd: f3 0f 58 08      addss  (%rax),%xmm1
c1: f3 0f 11 08      movss  %xmm1, (%rax)
for (int i=0; i<n; i++)
c5: 7e 41         jle    108 <saxpy+0x108>
    Y[i] = alpha*X[i] + Y[i];
c7: 4d 63 c0         movslq %r8d,%r8
ca: f3 42 0f 59 04 86 mulss  (%rsi,%r8,4),%xmm0
d0: 4a 8d 04 82      lea    (%rdx,%r8,4),%rax

```

```

d4:  f3 0f 58 00      addss  (%rax),%xmm0
d8:  f3 0f 11 00      movss  %xmm0,(%rax)
dc:  c3               retq
dd:  0f 1f 00        nopl   (%rax)
for (int i=0; i<n; i++)
e0:  31 c0            xor    %eax,%eax
e2:  66 0f 1f 44 00 00  nopw   0x0(%rax,%rax,1)
    Y[i] = alpha*X[i] + Y[i];
e8:  f3 0f 10 0c 86    movss  (%rsi,%rax,4),%xmm1
ed:  f3 0f 59 c8       mulss  %xmm0,%xmm1
f1:  f3 0f 58 0c 82    addss  (%rdx,%rax,4),%xmm1
f6:  f3 0f 11 0c 82    movss  %xmm1,(%rdx,%rax,4)
fb:  48 83 c0 01       add    $0x1,%rax
for (int i=0; i<n; i++)
ff:  39 c7            cmp    %eax,%edi
101: 7f e5            jg     e8 <saxpy+0xe8>
103: f3 c3            repz  retq
105: 0f 1f 00        nopl   (%rax)
108: f3 c3            repz  retq

```

The O3 optimization level, which is the highest one, performs the serious optimizing operation. In the above assembly code, the loop has been unrolled, in addition to that, function inlining and automatic vectorization have been performed.

c. Compile the saxpy procedure with -O2 optimization and the additional flag -mfpmath=387

The assembly code are generated by the following command: `gcc -c -g saxpy.c -o saxpyo2mfpm.o -O2 -mfpmath=387 --std=c99 objdump -d -S saxpyo2mfpm.o`

Here is the assembly code generated:

```

saxpyo2mfpm.o:      file format elf64-x86-64

Disassembly of section .text:

0000000000000000 <saxpy>:
/*
   Single-precision A*X plus Y, Y = alpha*X+Y
*/

```

```

void saxpy(int n, float alpha, float *X, float *Y) {
    for (int i=0; i<n; i++)
        0: 85 ff                test    %edi,%edi
        2: 7e 21                jle     25 <saxpy+0x25>
        4: f3 0f 11 44 24 f0    movss   %xmm0,-0x10(%rsp)
        a: 31 c0                xor     %eax,%eax
        c: d9 44 24 f0          flds    -0x10(%rsp)
        Y[i] = alpha*X[i] + Y[i];
    10: d9 c0                fld     %st(0)
    12: d8 0c 86                fmulss (%rsi,%rax,4)
    15: d8 04 82                faddss (%rdx,%rax,4)
    18: d9 1c 82                fstps  (%rdx,%rax,4)
    1b: 48 83 c0 01              add     $0x1,%rax
    for (int i=0; i<n; i++)
    1f: 39 c7                cmp     %eax,%edi
    21: 7f ed                jg      10 <saxpy+0x10>
    23: dd d8                fstp    %st(0)
    25: f3 c3                repz   retq

```

The loop structure is similar to the previous case in question a. However, the mathematical operation is different from the above. With `-mfpmath=387`, the assembly code produced to look much cleaner and seem to be more efficient.

### Question N°2

Compile the function `med3` with optimization level `-O2` and generate assembly code for it. Analyze and explain the assembly code like in question 1a

The assembly code are generated by the following command: `gcc -c -g med3.c -o med3o2.o -O2 --std=c99 objdump -d -S med3o2.o`

Here is the assembly code generated:

```

med3o2.o:      file format elf64-x86-64

Disassembly of section .text:

0000000000000000 <med3>:

int med3(int a, int b, int c) {
    int mid;
    if ( a > b )    {

```

```

0:  39 f7                cmp    %esi,%edi
int med3(int a, int b, int c) {
2:  89 f0                mov    %esi,%eax
if ( a > b ) {
4:  7e 12                jle    18 <med3+0x18>
    if ( c > b ) {
6:  39 d6                cmp    %edx,%esi
8:  7d 19                jge    23 <med3+0x23>
a:  39 d7                cmp    %edx,%edi
c:  89 d0                mov    %edx,%eax
e:  0f 4e c7             cmovle %edi,%eax
11: c3                  retq
12: 66 0f 1f 44 00 00     nopw   0x0(%rax,%rax,1)
    else {
        mid = b;
    }
}
else {
    if ( b > c ) {
18: 39 d6                cmp    %edx,%esi
1a: 7e 07                jle    23 <med3+0x23>
1c: 39 d7                cmp    %edx,%edi
1e: 89 d0                mov    %edx,%eax
20: 0f 4d c7             cmovge %edi,%eax
        else {
            mid = b;
        }
    }
    return mid ;
}
23:  f3 c3                repz retq

```

- how are the parameters passed to the procedure? The first parameter (a) is passed into register EDI and the second (b) in ESI, the third parameter (c) is passed into register EDX

- how is the condition block implemented? • how are the arithmetic operations done?

```
med3o2.o:      file format elf64-x86-64
```

```
Disassembly of section .text:
```

```

0000000000000000 <med3>:

int med3(int a, int b, int c) {
    int mid;
    if ( a > b )    {
        0:  39 f7                                cmp     %esi,%edi            #
compare b with a
int med3(int a, int b, int c) {
        2:  89 f0                                mov     %esi,%eax           #
move esi (b) to eax
        if ( a > b )    {
            4:  7e 12                                jle     18 <med3+0x18>       # if
less than or equal, jump to 18
            if ( c > b ) {
                6:  39 d6                                cmp     %edx,%esi           #
compare c with b
                8:  7d 19                                jge     23 <med3+0x23>       # if
greater than or equal
                a:  39 d7                                cmp     %edx,%edi           #
compare c with a
                c:  89 d0                                mov     %edx,%eax           #
move edx (c) to eax
                e:  0f 4e c7                                cmovle  %edi,%eax           #
move edi (a) to eax if less than or equal
                11:  c3                                retq                                     #
return
                12:  66 0f 1f 44 00 00                nopw    0x0(%rax,%rax,1)     # has
no effect, insert by the compiler to align branch targets
                mid = b;
            }
        }
    else {
        if ( b > c ) {
            18:  39 d6                                cmp     %edx,%esi           #
compare c to b
            1a:  7e 07                                jle     23 <med3+0x23>       # if
less than or equal, jump to 23 (return)
            1c:  39 d7                                cmp     %edx,%edi           #
compare c to a
            1e:  89 d0                                mov     %edx,%eax           #
move c to eax
            20:  0f 4d c7                                cmovge  %edi,%eax           #
move a to eax if greater than or equal
        else {
            mid = b;

```



```

    }
  }
  return mid ;
}
23:  f3 c3                      repz retq          #
return

```

• can you recognize any instructions that have no effect, i.e. that are inserted by the compiler to align branch targets? assembly

```
12:  66 0f 1f 44 00 00          nopw    0x0(%rax,%rax,1)    # has
no effect, insert by the compiler to align branch targets
```

### Question N°3

The procedure in the attached file call\_max.c illustrates how procedure calls are implemented in the generated assembly language

a. The assembly code are generated by the following command: gcc -c -g call\_max.c -o call\_max-fno-inline.o -O2 -fno-inline --std=c99 objdump -d -S call\_max-fno-inline.o

Here is the assembly code generated with explanations:

```

call_max-fno-inline.o:      file format elf64-x86-64

Disassembly of section .text:

0000000000000000 <max>:
int max(int a, int b) {
/*
The first parameter (a) is passed into register EDI and the second
(b) in ESI
*/
    0: 39 f7                cmp     %esi,%edi
# compare b to a
    2: 89 f0                mov     %esi,%eax
# move b to eax
    4: 0f 4d c7             cmovge %edi,%eax
# move a to eax if greater than or equal
    if (a>b) return a;
    else return b;
}
    7: c3                  retq

```

```

# return
    8: 0f 1f 84 00 00 00 00  nopl    0x0(%rax,%rax,1)
# has no effect
    f: 00

0000000000000010 <select>:

void select (int *a, int *b, int len) {
/*
The first parameter (*a) is passed into register EDI and the second
(*b) in ESI, the third parameter (len) is passed into register EDX
*/
    for (int i=0; i<len; i++) {
10: 85 d2                      test    %edx,%edx
# if len = 0
12: 7e 3c                      jle     50 <select+0x40>
# then return by jumping into 50
void select (int *a, int *b, int len) {
14: 41 55                      push    %r13
# push r13
16: 8d 42 ff                  lea     -0x1(%rdx),%eax
# eax = len - 1
19: 41 54                      push    %r12
# push r12
1b: 4c 8d 2c 85 04 00 00  lea     0x4(,%rax,4),%r13
# r13 = r12 + 4
22: 00
23: 49 89 f4                  mov     %rsi,%r12
# move rsi to r12
26: 55                      push    %rbp
# push register base pointer
27: 48 89 fd                  mov     %rdi,%rbp
# move rdi to rbp
2a: 53                      push    %rbx
# push rbx
    for (int i=0; i<len; i++) {
2b: 31 db                      xor     %ebx,%ebx
# i = 0
2d: 0f 1f 00                  nopl    (%rax)
# no effect
    a[i] = max(a[i], b[i]);
30: 41 8b 34 1c              mov     (%r12,%rbx,1),%esi
# move r12 to esi
34: 8b 7c 1d 00              mov     0x0(%rbp,%rbx,1),%edi
# move rbp to edi

```

```

    38: e8 00 00 00 00      callq 3d <select+0x2d>
# calq 3d
    3d: 89 44 1d 00      mov    %eax,0x0(%rbp,%rbx,1)
# move eax to rbp
    41: 48 83 c3 04      add    $0x4,%rbx
# rbx += 4
    for (int i=0; i<len; i++) {
    45: 4c 39 eb          cmp    %r13,%rbx
# compare r13 with rbx
    48: 75 e6            jne    30 <select+0x20>
# if not equal, jump to 30
    }

}

    4a: 5b              pop    %rbx
# pop rbx
    4b: 5d              pop    %rbp
# pop rbp
    4c: 41 5c          pop    %r12
# pop r12
    4e: 41 5d          pop    %r13
# pop r13
    50: f3 c3          repz retq
# return

```

b. The assembly code are generated by the following command: `gcc -c -g call_max.c -o call_max.o -O2 --std=c99 objdump -d -S call_max.o`

Here is the assembly code generated:

```

call_max-fno-inline.o:      file format elf64-x86-64

Disassembly of section .text:

0000000000000000 <max>:
int max(int a, int b) {
/*
The first parameter (a) is passed into register EDI and the second
(b) in ESI
*/
    0:  39 f7              cmp    %esi,%edi          #
compare b to a

```

```

    2:  89 f0                mov     %esi,%eax                #
move b to eax
    4:  0f 4d c7              cmovge %edi,%eax                #
move a to eax if greater than or equal
    if (a>b) return a;
    else return b;
}
    7:  c3                      retq                     #
return
    8:  0f 1f 84 00 00 00 00    nopl     0x0(%rax,%rax,1)        # no
effect
    f:  00

00000000000000010 <select>:

void select (int *a, int *b, int len) {
/*
The first parameter (*a) is passed into register EDI and the second
(*b) in ESI, the third parameter (len) is passed into register EDX
*/
    for (int i=0; i<len; i++) {
    10:  31 c0                  xor     %eax,%eax                # set
i = 0
    12:  85 d2                  test    %edx,%edx                # if
len = 0
    14:  7e 1f                  jle     35 <select+0x25>          #
then return by jumping to 35
    16:  66 2e 0f 1f 84 00 00    nopw    %cs:0x0(%rax,%rax,1)    # no
effect, used to align
    1d:  00 00 00
    20:  8b 0c 86              mov     (%rsi,%rax,4),%ecx        #
move rsi to ecx
    23:  39 0c 87              cmp     %ecx,(%rdi,%rax,4)        #
compare ecx to rdi
    26:  0f 4d 0c 87          cmovge (%rdi,%rax,4),%ecx        #
move rdi to ecx if greater than or equal
    a[i] = max(a[i], b[i]);
    2a:  89 0c 87              mov     %ecx,(%rdi,%rax,4)        #
move ecx to rdi
    2d:  48 83 c0 01          add     $0x1,%rax                # i++
    for (int i=0; i<len; i++) {
    31:  39 c2                  cmp     %eax,%edx                #
compare i to len
    33:  7f eb                  jg      20 <select+0x10>          #
jump to 20 if greater than

```

```
35:  f3 c3          repz retq          #  
return
```

With function inlining enabled, the code is more simple and uses fewer registers to perform.